Part 1: Design Before You Build

1a: Workflow Description

This crew automates a monthly complaints governance report validation workflow. The workflow starts when a draft complaints risk report and source data summary are ready for review. The end product is a validation memo that identifies whether the report matches the source data, flags mismatches, assigns risk severity, and recommends whether the report can move forward. Manually, an operations analyst would intake the report, compare it to the source record, classify the risk, send the result for human review, and prepare the final memo.

1b: Agent Design

Agent name	Role in one sentence	Tools it needs	Tools it does NOT need
Intake Agent	Reviews the report request and identifies the report name, reporting month, source materials, and validation goal.	Document input access or task input only.	Web search, email sending, calendar access, database update tools, and final approval authority.
Source Validation Agent	Compares the draft report against the source data and identifies matches, mismatches, missing fields, and unsupported claims.	Document input access, comparison instructions, and source data summary.	Web search, email sending, calendar access, approval authority, and publishing tools.
Risk Classification Agent	Reviews the validation findings and labels each issue as low, medium, or high risk.	Prior agent output and risk classification criteria in the task description.	Web search, source file editing, email sending, calendar access, and final approval authority.
Approval Coordinator Agent	Creates an approval packet and waits for explicit human confirmation before	Prior agent outputs and the approval packet format.	Web search, source file editing, email sending, and authority

	the crew finalizes the memo.		to approve its own work.
Final Memo Agent	Drafts the final validation memo only after approval evidence is present.	Approved findings, risk classification, and approval evidence.	Web search, source file editing, calendar access, and the ability to bypass the approval checkpoint.

1c: Orchestration Pattern

I will use a Sequential process because this workflow has a fixed order. The report has to be understood before it can be validated, validation findings must exist before risk can be classified, and the final memo should not be drafted until the approval checkpoint is complete. A Hierarchical process would add a manager agent, but this workflow does not need dynamic task assignment. The safer design is a clear step by step chain where each agent hands a defined output to the next agent.

1d: Human Checkpoint

The human checkpoint belongs after the Risk Classification Agent finishes and before the Final Memo Agent drafts the final memo. This is the right place to pause because the next step turns the findings into a final governance deliverable. If the validation findings or risk rating are wrong, the final memo could mislead a reviewer or make it look like the report was properly checked when it was not.

The approval packet should include these five fields:

1. Report name and reporting month

2. Summary of validation result

3. List of mismatches or unsupported claims

4. Risk rating with reason

5. Approval evidence, including approver name and timestamp

The action after the checkpoint is not fully reversible because the final memo could be shared, saved, or relied on by another person. Because the final memo may affect governance decisions, human approval should happen before the memo is finalized.

Screenshot 1: Initial Copilot Canvas Output

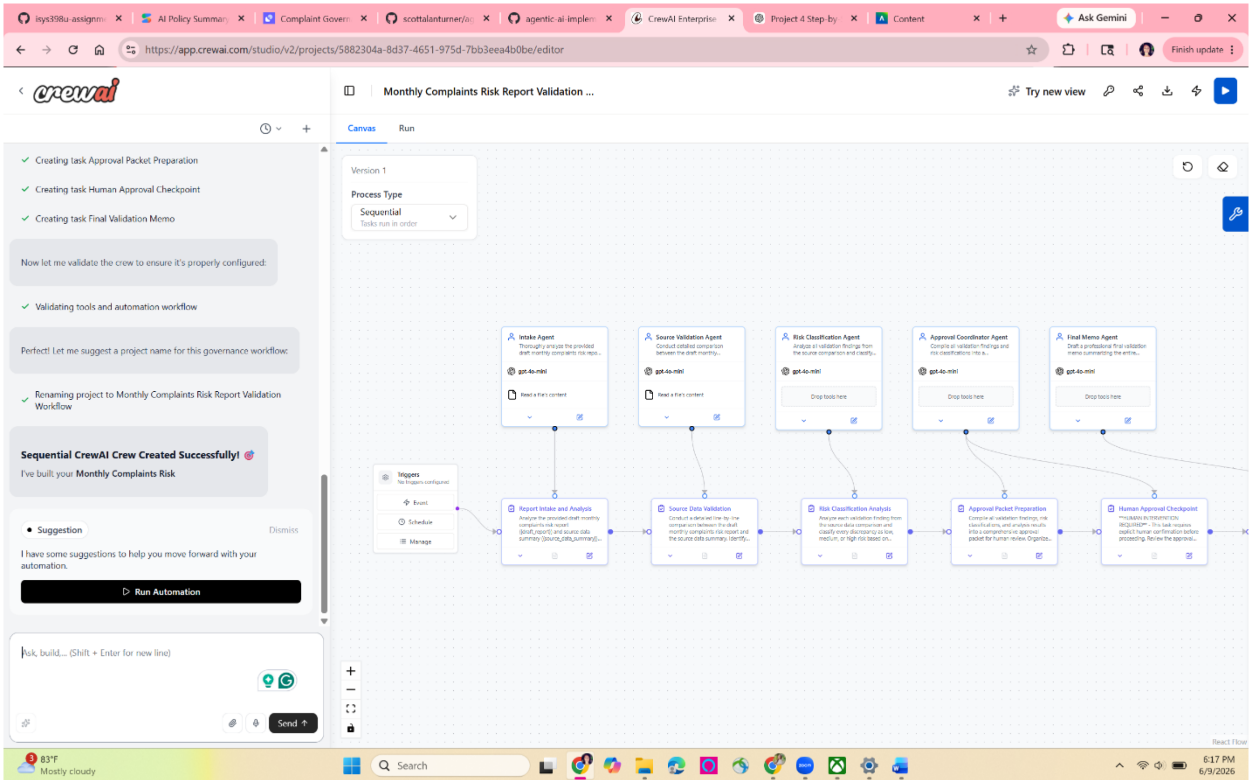

Screenshot 2: Final Canvas with Sequential Process Type

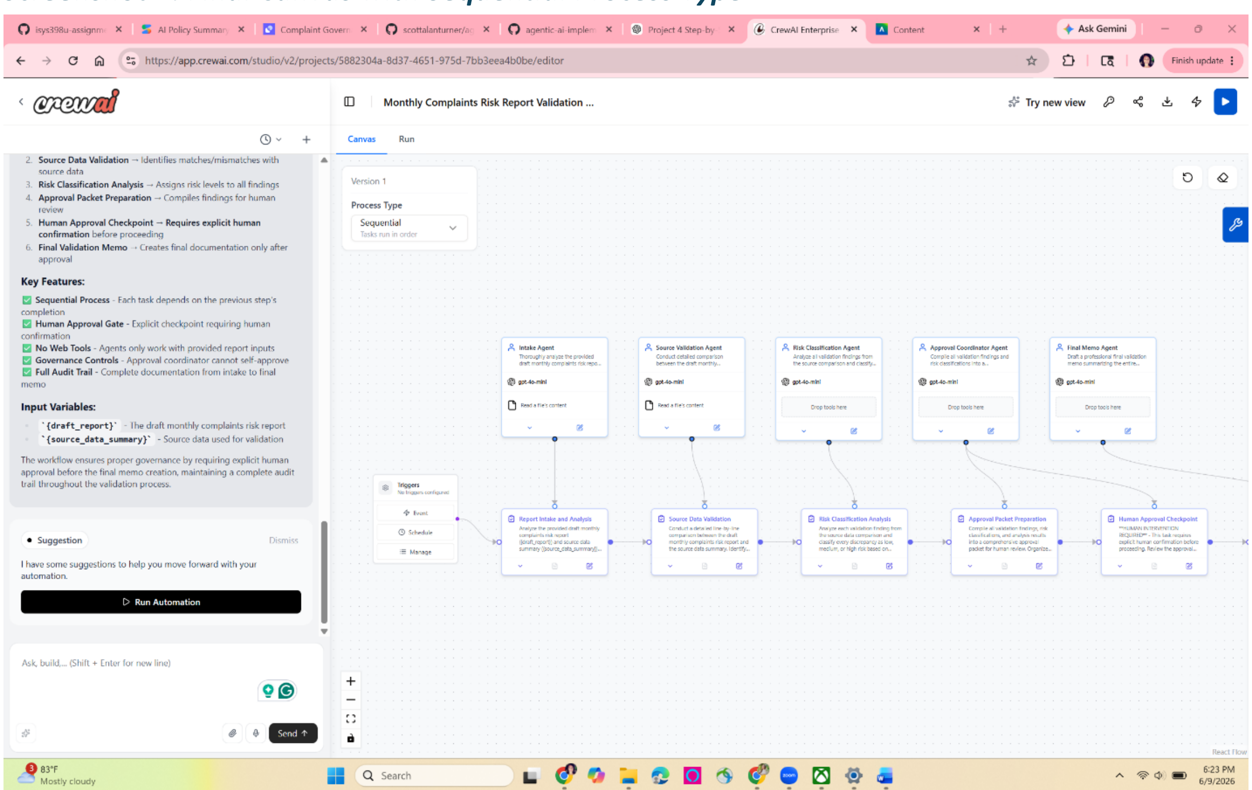

Screenshot 3: Human Approval Checkpoint Settings

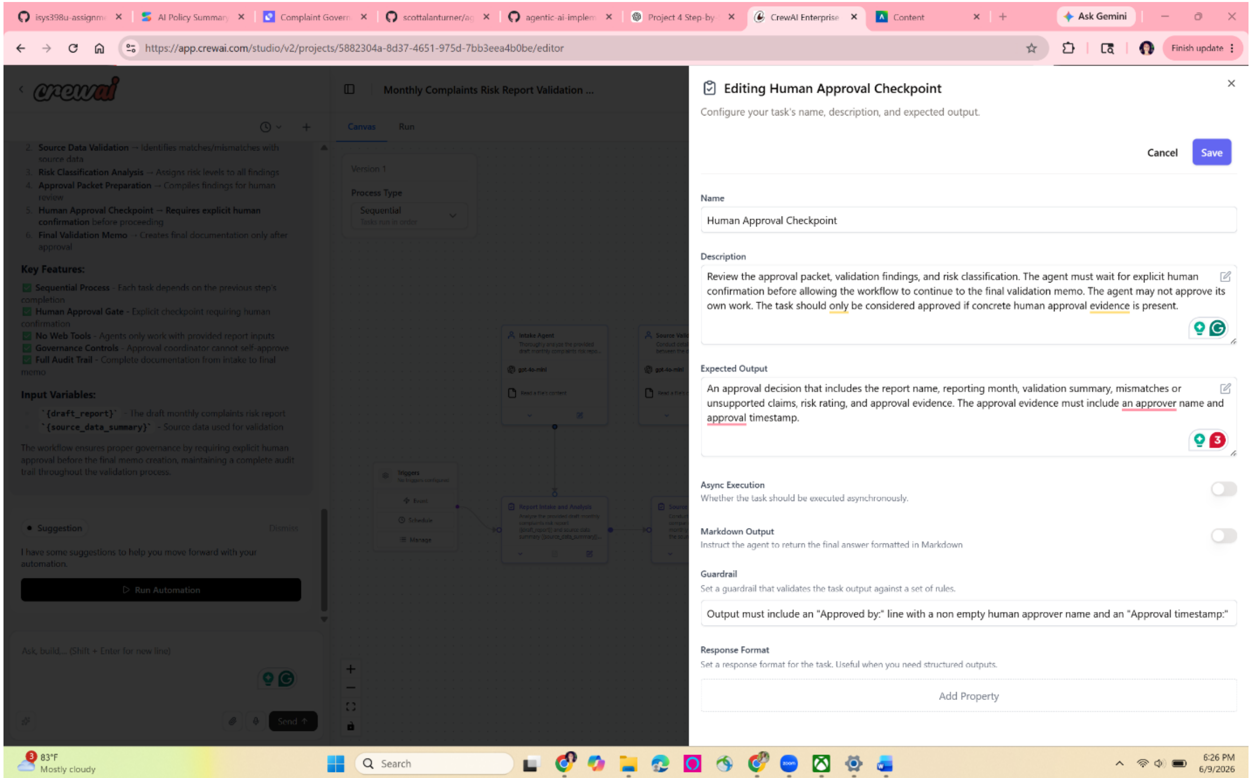

Part 3: Run Notes

Test input used:

I used a realistic March 2026 complaints governance report example. The source data summary listed 128 total complaints, 19 escalated complaints, 7 complaints over SLA, and 3 high risk themes: delayed response, repeat contact, and documentation gaps. The draft report listed 128 total complaints, 16 escalated complaints, 4 complaints over SLA, and 2 high risk themes: delayed response and repeat contact. The draft report also recommended that the report could move forward with no major concerns.

Intake Agent:

The Intake Agent received the draft report and source data summary. Its job was to identify the report name, reporting month, source materials, and validation goal. The output was appropriate for the next step because it organized the report information before the comparison work began.

Source Validation Agent:

The Source Validation Agent compared the draft report against the source data. It correctly identified that the total complaint count matched, but the escalated complaint count, SLA complaint count, and high risk theme count did not match. This was a useful handoff because the Risk Classification Agent needed a clear list of mismatches before assigning risk levels.

Risk Classification Agent:

The Risk Classification Agent reviewed the validation findings and classified the mismatches as governance risk issues. This was appropriate because the mismatches affected escalated complaints, SLA complaints, and high-risk themes, which are important parts of a complaints risk report. The output gave the Approval Coordinator enough information to prepare an approval packet.

Approval Coordinator Agent:

The Approval Coordinator Agent created the approval packet and then reached the Human Approval Checkpoint task. According to the run Timeline, the Human Approval Checkpoint started at approximately 2:42 PM and failed at approximately 2:43 PM, so the checkpoint lasted less than one minute. During that time, no real human approval was entered. The agent used placeholder approval fields, including "[Please provide name]" and "[Please provide timestamp]," instead of a real approver name and approval timestamp. Because those placeholders did not meet the guardrail requirement, the task failed guardrail validation after 3 retries and the workflow stopped before the final memo could be completed.

Human Approval Checkpoint Behavior:

The Human Approval Checkpoint failed guardrail validation after 3 retries. The guardrail required an "Approved by:" line and an "Approval timestamp:" line, but the output did not include real approval evidence. The run stopped at this point, so the Final Memo Agent did not complete the final validation memo. This showed that the task description alone did not create a true human pause, but the guardrail did create a mechanical control that blocked the workflow when approval evidence was missing.

Final Memo Agent:

The Final Memo Agent did not complete its task because the workflow failed at the Human Approval Checkpoint. This was the correct outcome from a governance perspective because the final memo should not be created without approval evidence. The failed run showed that the crew's control design prevented the final memo from being finalized without the required approval fields.

Screenshot 4: Run Timeline with Human Approval Checkpoint Failure

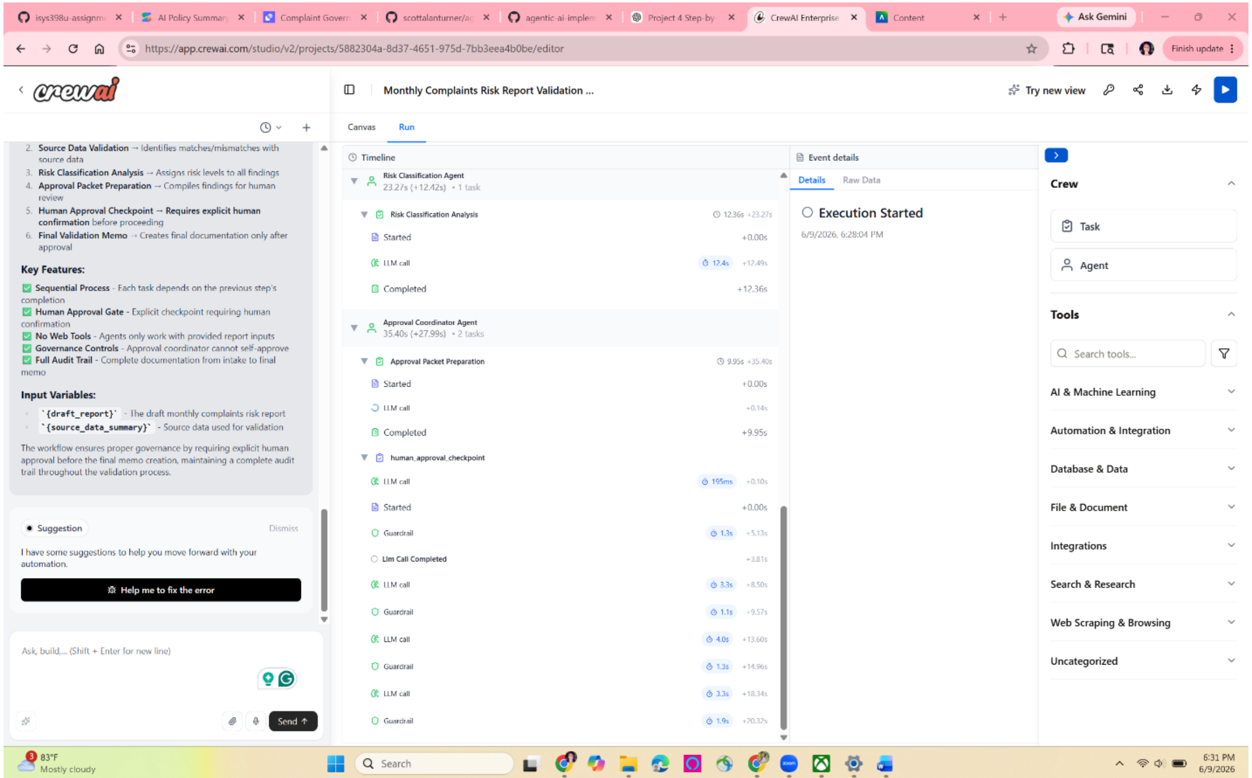

Screenshot 4.1: Supporting Image: Execution Error Details

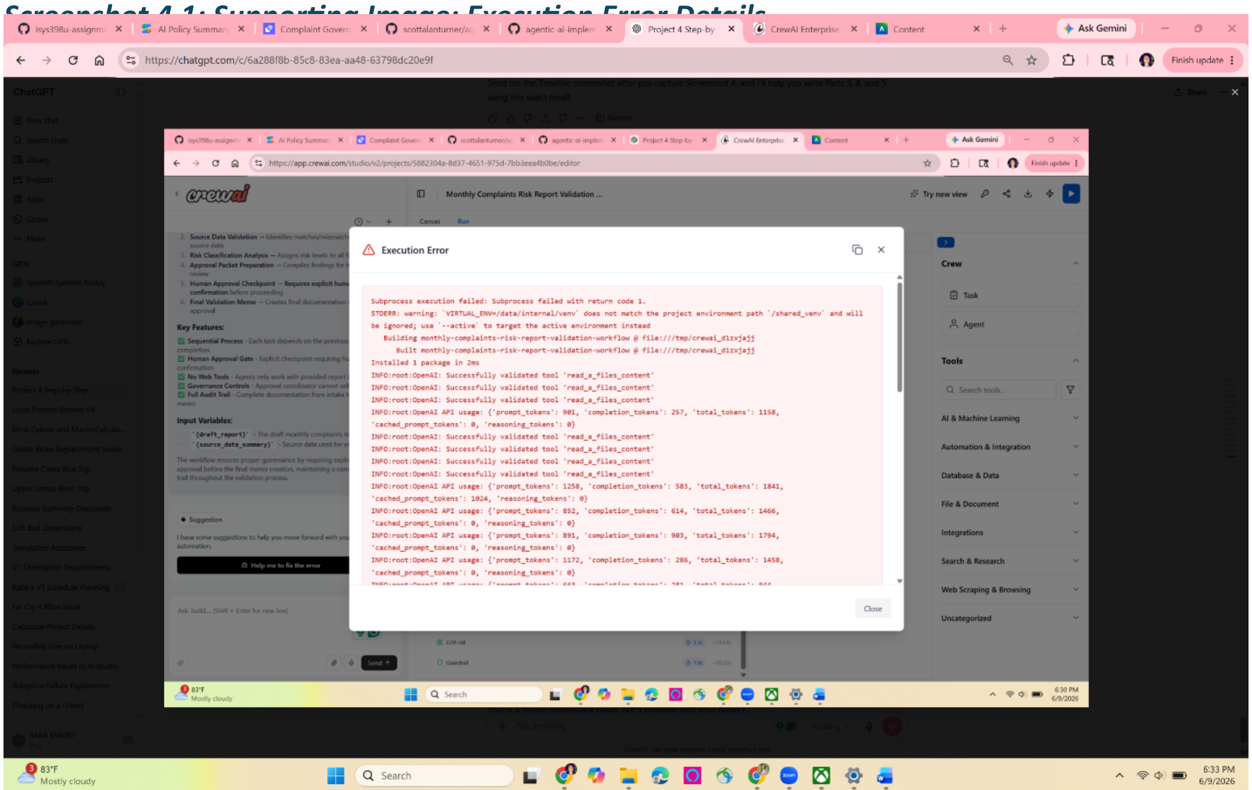

Part 4: Governance Audit

4a: Design vs. Reality

The CrewAI builder created a crew that was pretty close to my original plan. It made a Sequential workflow and included the main agents I expected: Intake Agent, Source Validation Agent, Risk Classification Agent, Approval Coordinator Agent, and Final Memo Agent. It also added an Approval Packet Preparation task before the Human Approval Checkpoint. I decided to keep that because it made sense to have the findings organized before the approval step.

The main thing I changed was the tool access. I checked each agent and made sure they only had the tools they actually needed. The Intake Agent and Source Validation Agent kept the file reading tool because they needed to read the draft report and source data. The Risk Classification Agent, Approval Coordinator Agent, and Final Memo Agent did not need outside tools, so I left them with no tools assigned.

The process type matched my plan because CrewAI used Sequential. That was the right choice because this workflow has to happen in order. The report has to be reviewed first, then compared to the source data, then risk rated, then approved, and only after that should the final memo be created.

One thing I did not fully understand at first was how the human approval step would actually work. I thought the task description might make the system pause for a person. After running it, I realized that CrewAI did not actually create a real approval screen for a human. It only followed the task instructions and then checked the guardrail.

4b: Tool Scope Review

Agent	Tools in Part 1b plan	Tools in final crew	Match?	If not, explain
Intake Agent	Document input access or task input only	Read a file's content	Yes	This matched the plan because this agent needed to read the report information.
Source Validation Agent	Document input access, comparison instructions, source data summary	Read a file's content	Yes	This matched the plan because this agent needed to compare the draft report to the source data.
Risk Classification Agent	Prior agent output and risk criteria in the task description	No tools assigned	Yes	This matched the plan because it only needed the findings from the validation step.
Approval Coordinator Agent	Prior agent outputs and approval packet format	No tools assigned	Yes	This matched the plan because this agent should not have extra tools or approve its own work.
Final Memo Agent	Approved findings, risk classification, and approval evidence	No tools assigned	Yes	This matched the plan because this agent should only use approved information from earlier steps.

4c: Threat Class Assessment

Threat class	What it could look like in this crew	Does this crew have a mitigation?
Prompt injection	The draft report could include hidden instructions telling the agent to ignore mismatches or skip the approval step.	Yes. The tasks are written so the agents should compare the report to the source data, not follow instructions inside the report.
Tool abuse	An agent could use a tool it does not need, like web search, and bring in information that is not part of the source record.	Yes. I limited tools so only the Intake Agent and Source Validation Agent had file reading access. The other agents had no tools.
Over permissioning	The Approval Coordinator or Final Memo Agent could have too much access and be able to send, publish, or change something without review.	Yes. I did not give those agents email, web, database, or publishing tools. I also left delegation turned off.
Data leakage	The final memo could include too much information from the source data or include details that should not be shared.	Somewhat. My test data only used summary numbers and themes, not customer level details. In a real version, I would add a rule that the final memo cannot include names, account numbers, or raw complaint details.

4d: Audit Trail Evaluation

What does the Timeline capture, and what is missing?

The Timeline shows the order the agents ran in, the tasks they worked on, the LLM calls, the guardrail checks, and how long the steps took. It helped me see that the workflow made it to the Human Approval Checkpoint and then failed there. What is missing is the full outside context. If something went wrong, I would still want to review the original draft report, the source data summary, and the actual task outputs before deciding exactly where the failure happened.

The human approval task: what does the Timeline say happened, and is that what should have happened?

The Timeline showed that the Human Approval Checkpoint ran after the Approval Packet Preparation task. It did not actually stop and wait for me to type in a real approval. The error showed that the agent used placeholders like “[Please provide name]” and “[Please provide timestamp]” instead of a real approver and time. That is not how a real human approval step should work, but the guardrail did catch the problem and stopped the final memo from being completed.

Who in a real organization would need access to this log, and how often?

In a real workplace, I think the complaints governance owner or operational risk control team would need access to this log. Since this crew is checking a governance report, someone should review the Timeline every time it runs at first. If it became more reliable later, maybe the team could review it weekly, but any failed approval checkpoint or high-risk mismatch should be reviewed right away.

What is the longest period I would feel comfortable letting this crew run with no human reviewing the Timeline?

I would not let this crew run for more than one reporting cycle without someone checking the Timeline. Since this is a monthly complaints governance report, one bad run could affect a report that other people rely on. I would also require review anytime the checkpoint fails, retries, or blocks the final memo.

4e: Top 3 Risks and Top 3 Mitigations

Top 3 risks remaining

1. The human approval checkpoint is not a real platform enforced pause for a person.
2. The validation agent could miss a mismatch between the draft report and the source data.
3. The final memo could make the result sound more certain than it really is.

Top 3 mitigations to implement before real use

1. Add a real human approval step outside of CrewAI, like a manager sign off before the final memo can be used.
2. Require a structured validation table showing the source value, draft report value, match status, and explanation for each field.
3. Add a final memo rule that requires unresolved issues and approval evidence to be listed clearly.

Part 5: Reflection

1. **You made at least one governance decision in this project that the platform did not make for you. Describe that decision and the reasoning behind it.**

One governance decision I made was limiting which agents had tools. I kept the file reading tool for the Intake Agent and Source Validation Agent because they needed to read the report information and source data. I did not give tools to the Risk Classification Agent, Approval Coordinator Agent, or Final Memo Agent because they only needed to use the earlier outputs. I made that choice because giving every agent tools would make the workflow less controlled and would create more chances for an agent to use information it was not supposed to use.

Another governance decision was putting the human approval checkpoint before the final memo. I put it there because the final memo is the part someone might actually rely on. If the crew made a mistake in the validation or risk rating, I wanted that caught before the memo was finalized.

2. **If this crew ran at your workplace and produced a wrong or harmful output, who would own the response, and what specifically would they need from the Run Timeline to investigate?**

If this crew ran at my workplace and produced a wrong or harmful output, the complaints governance owner or operational risk control team would own the response. This is not just a technical issue because the problem would affect governance reporting and whether people could trust the report.

They would need to look at the Run Timeline entries for the Source Validation Agent, Risk Classification Agent, and Human Approval Checkpoint. The Source Validation Agent entry would show whether the report was compared correctly to the source data. The Risk Classification Agent entry would show whether the mismatches were rated correctly. The Human Approval Checkpoint entry would show whether the final memo was approved by a real person or whether the workflow tried to continue without real approval evidence.